

KAGGLE COMPETITION

PREDICTING HOUSING PRICES

Team MealPals

Christian Opperman

Melanie Zheng

Paul Choi



Introduction

Challenge

- Predict home prices in Ames, Iowa using advanced machine learning techniques

Why?

- You are a Real Estate Investor with the opportunity to invest in a portfolio of 1,459 homes
- This deal has a specific asking price for the entire package of homes and contains underlying data with each house's features
- You need help deciding if the transaction is worth the asking price — in other words, whether it represents a good enough return on investment (“ROI”)
- *In this scenario, we assume **(1)** time is not a factor and **(2)** these homes will ultimately be re-sold at or near their predicted prices*



Solving the Problem

Approach:

- Build a machine learning model using historical housing data (the “train set”) to predict values on another group of homes (the “test set”)
- Use this information to determine a final investment recommendation on the purchase of this group of homes



Our Solution

Recommendation:

- Our best model for predicting the true values of the home portfolio came from a weighted average of our Stacked and Light GBM models

Model	Details	Kaggle Score
Weighted Average of Best Models	0.7 Stacked Model + 0.3 Light GBM	0.12873

- Our final recommendation is based on if the deal’s “asking” price is above or below our total valuation of **\$262 million** based on the aggregated predicted prices for all 1,459 homes
- We include some margin of error in our analysis and exclude any further ROI requirements for the purposes of this example

Exploratory Data Analysis




First Glance at the Kaggle Data

Two datasets provided

- Train set
 - 1,460 observations
 - 79 predictor variables (excluding 'SalePrice', 'Id')
 - Numeric variables: 28
 - Categorical variables: 51
- Test set
 - 1,459 observations, 79 predictor variables (excluding 'Id')

Exploratory Data Analysis

A quick look at our target variable: **SalePrice**

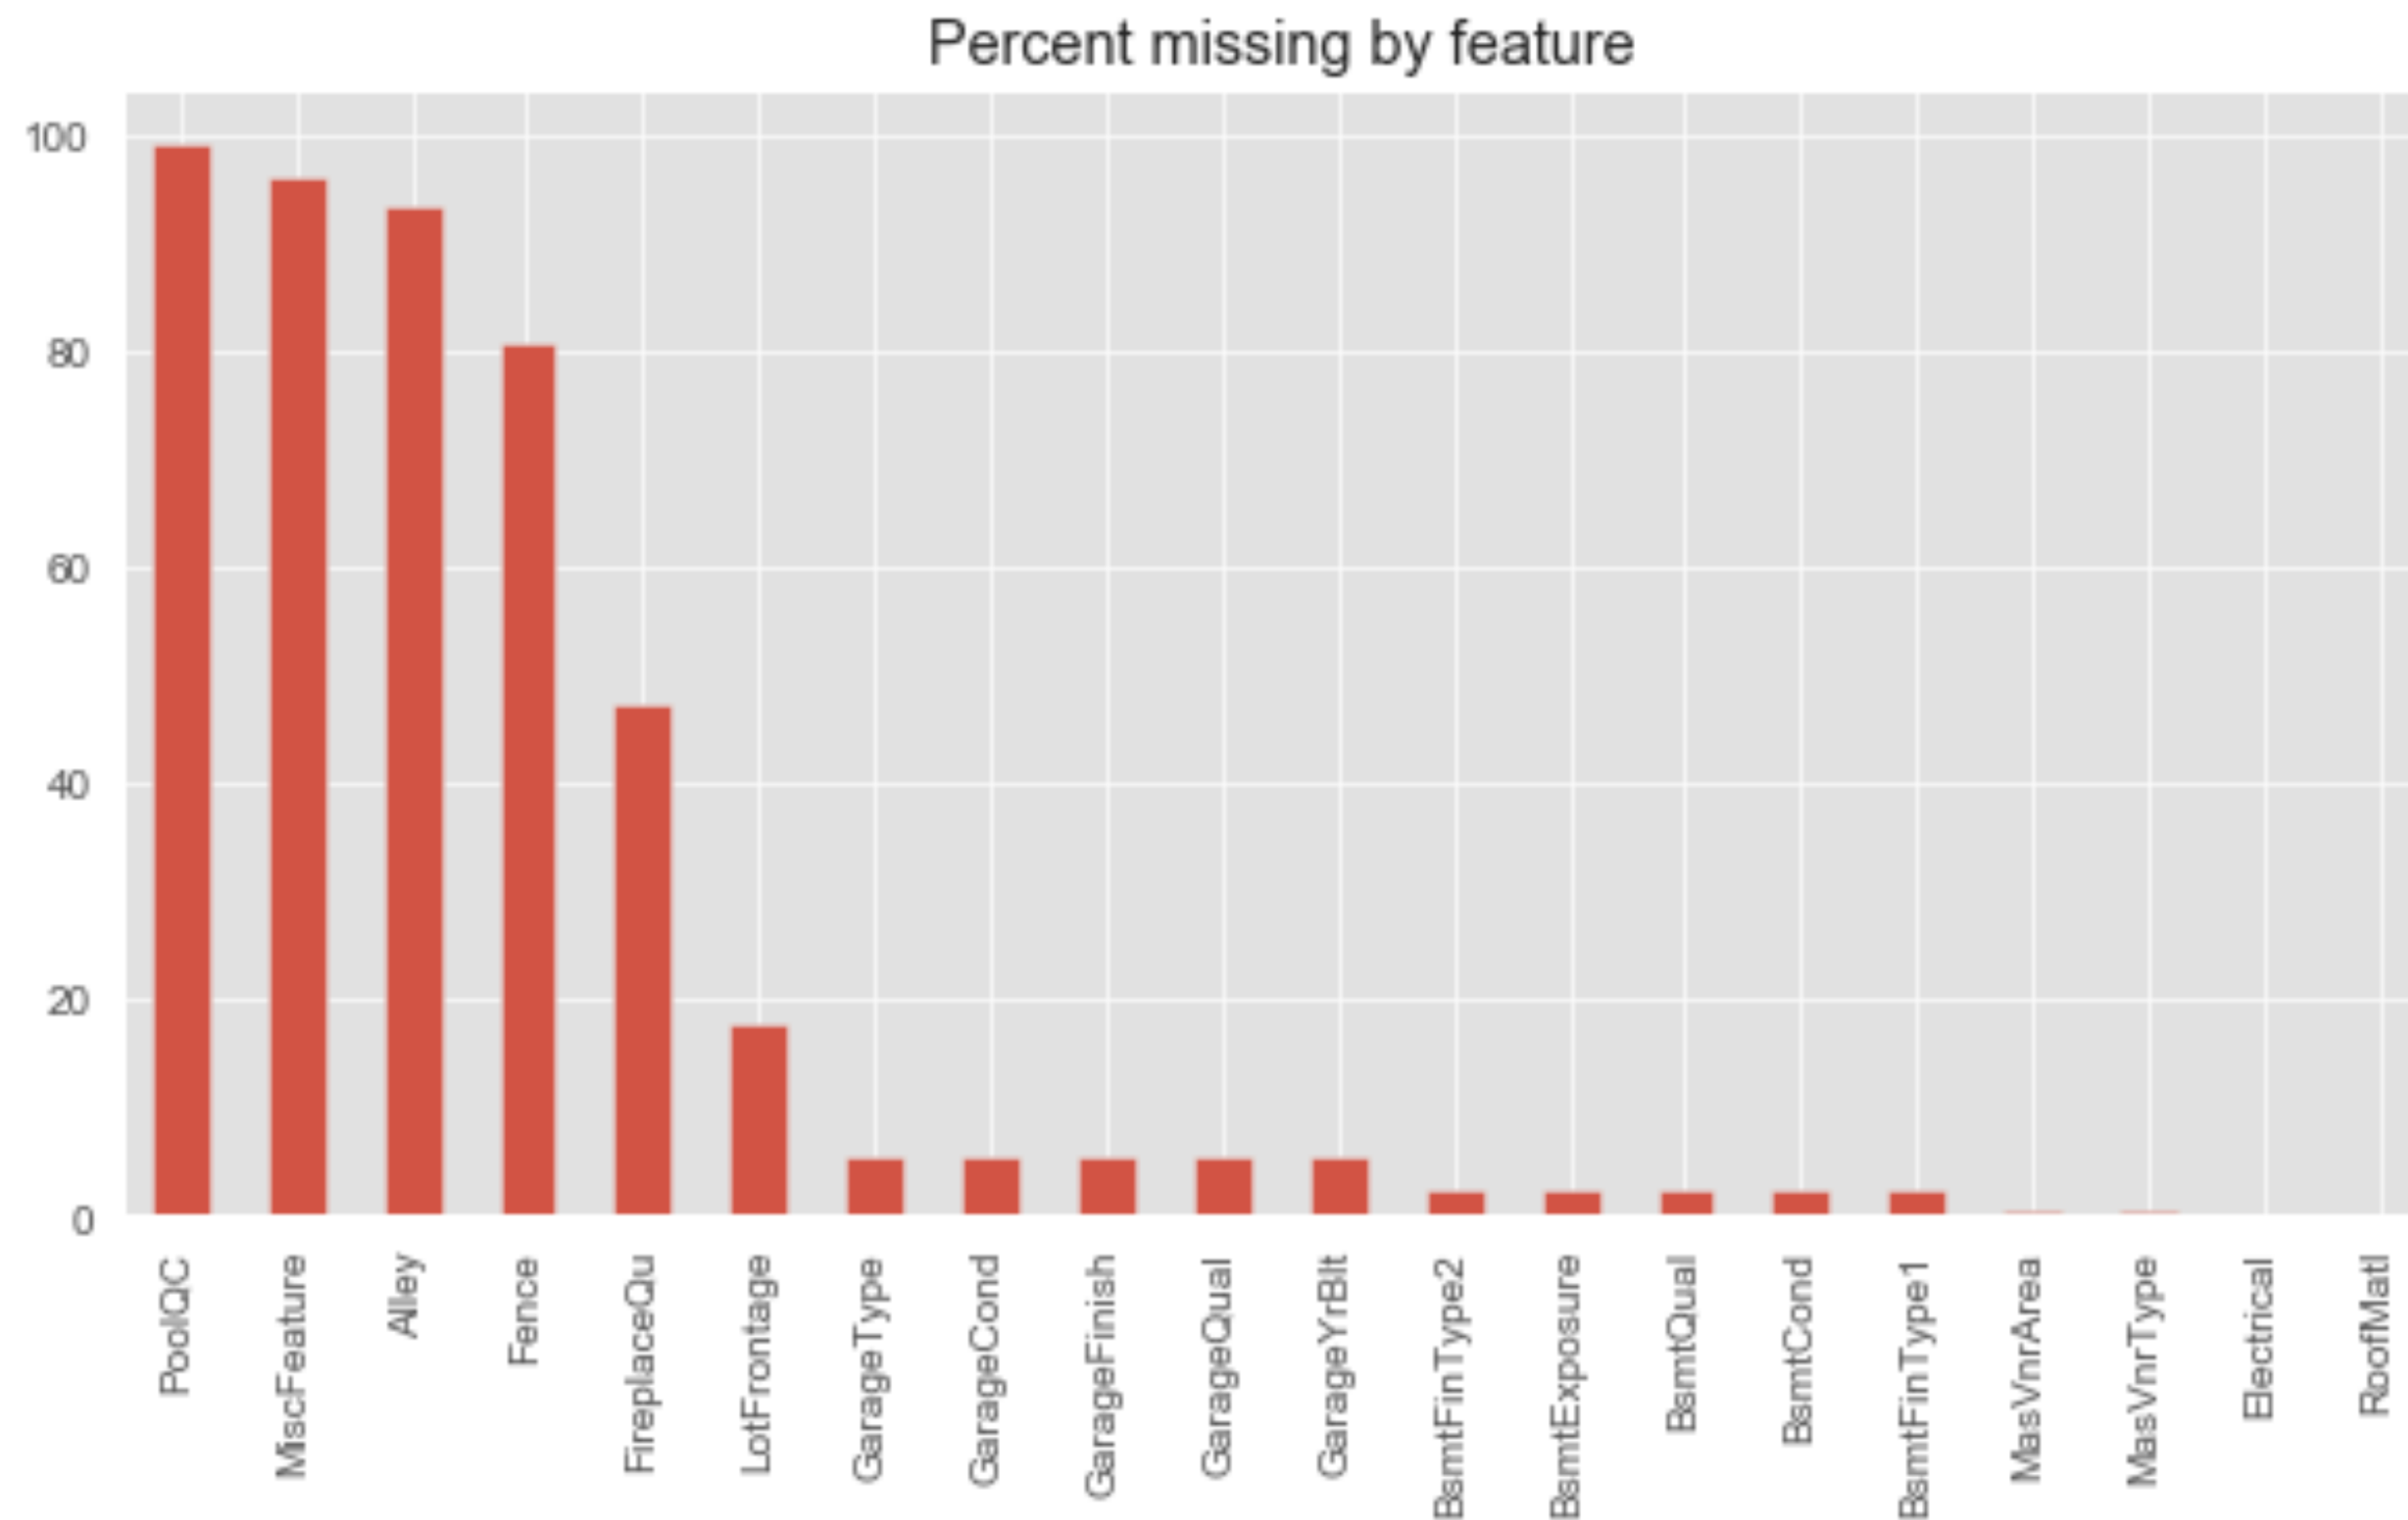


EDA: Examining Variable Distribution




EDA: Looking for Missingness

Check the proportion of missing data by variable



EDA: Identifying Outliers


Isolating and removing outliers



EDA: Examining Multicollinearity

Examples of variables with high correlation with each other:

- 1stFlrSF & TotalBsmtSF
- TotRmsAbvGrd & GrLivArea
- GarageArea & GarageCars
- YearBuilt & GarageYrBuilt



Data Preprocessing



Data Preprocessing Methodology

1. Remove outliers in training data
2. Impute data:
 - Impute pseudo-missing values
 - Impute true-missing values
 - Re-engineer categorical features as necessary
3. Add new feature variables
4. Dummify categorical feature variables
5. Remove redundant feature variables

Imputation: Pseudo vs. Actual Missingness

- There were a number of feature variables that contained missing values that did not, in fact, represent missing data.
- The general philosophy for a given variable X (where X represents a housing feature such as a pool, a fireplace, etc.) with this “pseudo-missingness” was to impute missing values as “No X ”

Pseudo Missing Values

Alley, BsmtCond, BsmtQual, BsmtFinType1,
Fence, Fireplace, GarageCond, GarageFinish,
GarageQual, GarageType, GarageYrBlt,
MasVnrType, MiscFeature, PoolQC

Actual Missing Values

Electrical,
MasVnrArea,
LotFrontage,
BsmtExposure,
BsmtFinType2



Feature Engineering

Generated New Features


- $\text{Total SF} = \text{1stFlrSF} + \text{2ndFlrSF} + \text{TotalBsmtSF} - \text{Numeric}$
- $\text{TotalFullBath} = \text{BsmtFullBath} + \text{FullBath} - \text{Numeric}$
- $\text{TotalHalfBath} = \text{BsmtHalfBath} + \text{HalfBath} - \text{Numeric}$
- IsPool — Categorical
- IsGarage — Categorical

Dummified Categorical Feature Variables

Removed Redundant Variables

SalePrice Transformation: Log vs. Box-Cox

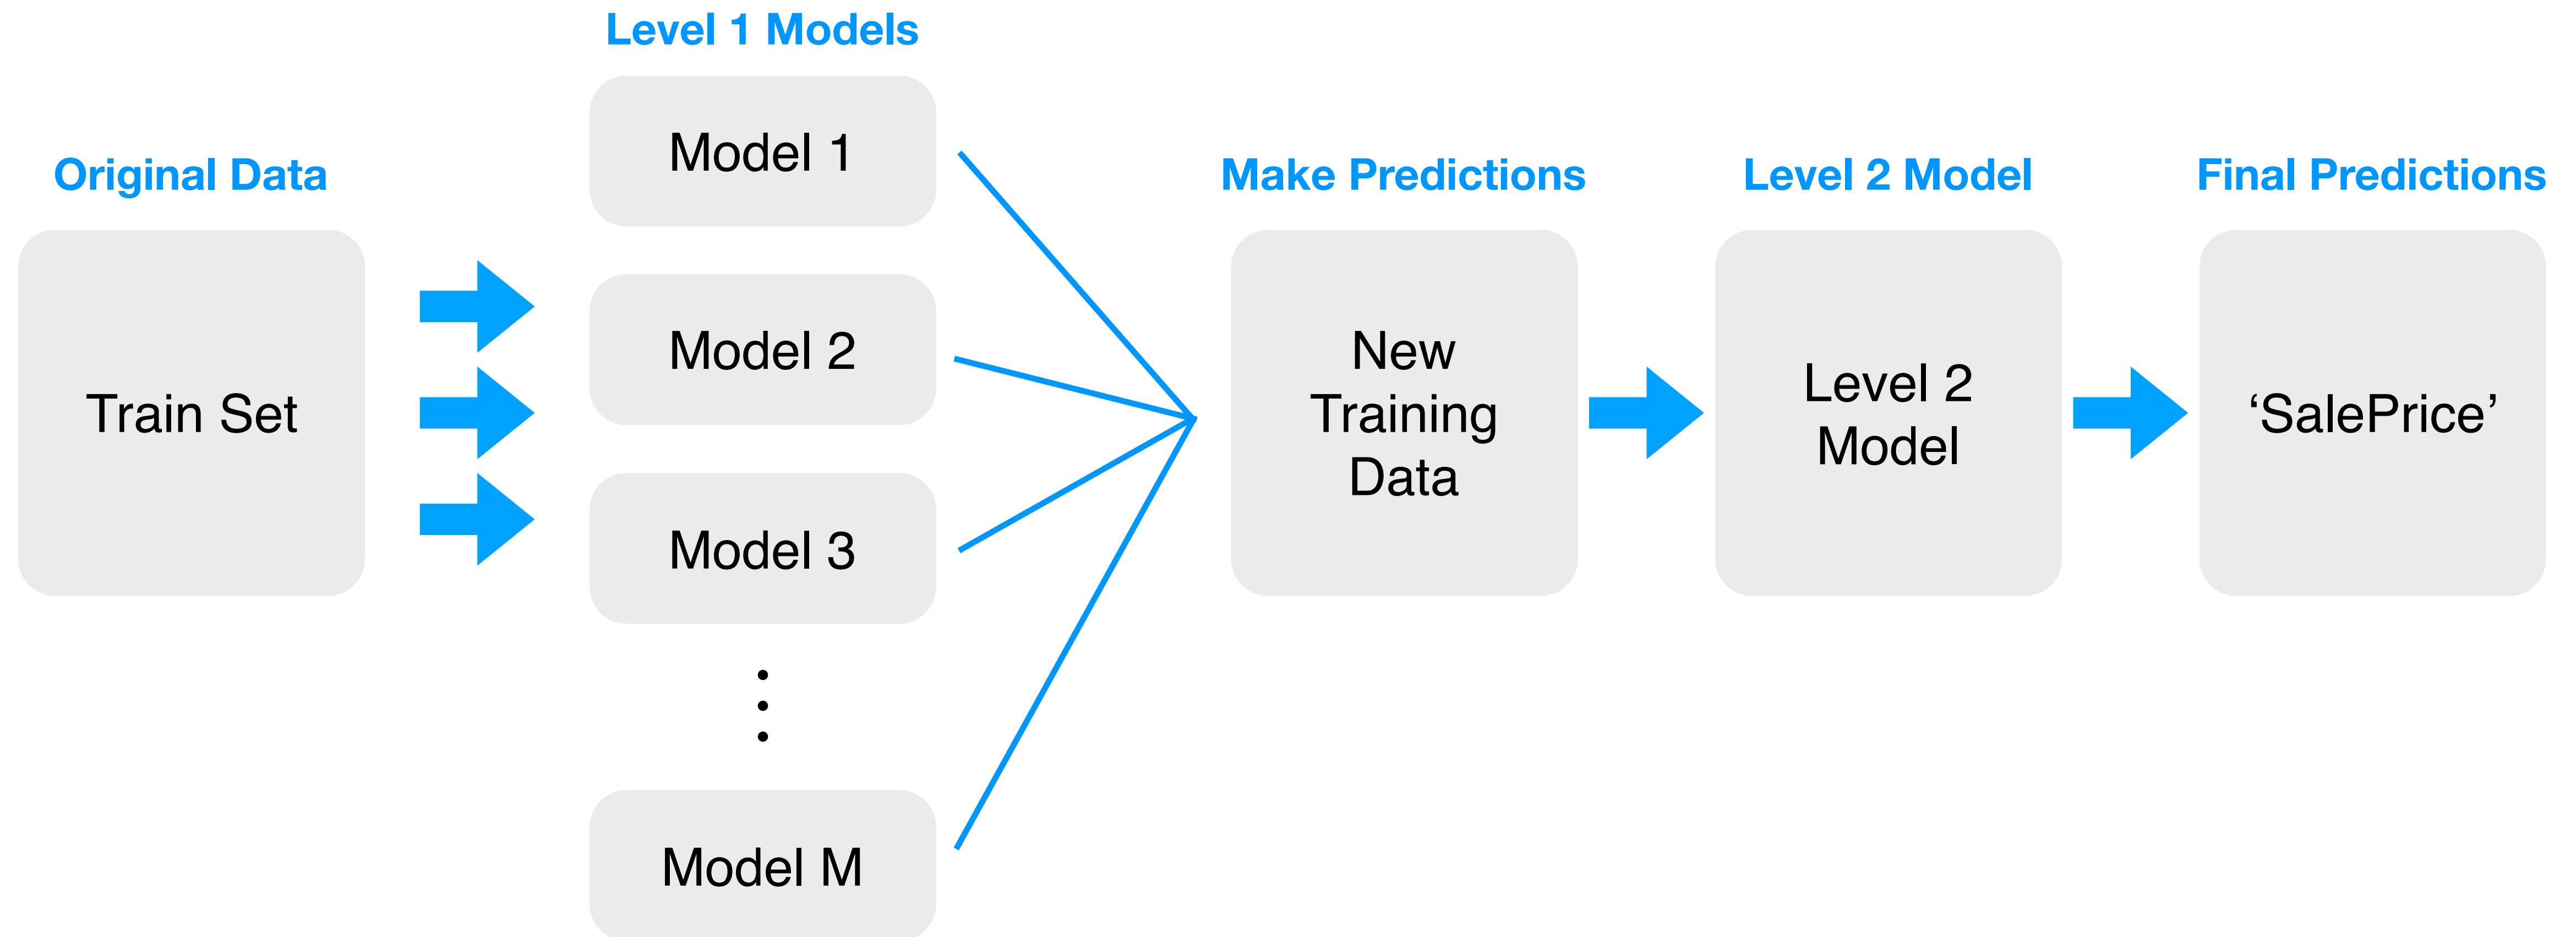
We decided to utilize a Box-Cox transformation over a Log transformation as it provided slightly better predictive power





Models



Modeling Methodology





Models Tested

Model Level	Model	Cross Validation Score (R^2)	Kaggle Score
Level I	Ridge	0.11739	0.14049
	Lasso	0.11413	0.13839
	Elastic Net	0.11800	0.13974
	Catboost	0.11678	0.13343
	Gradient Boost	0.12088	0.13435
	XGBoost	0.11952	0.13528
	LightGBM	0.11925	0.13123
Level II	Simple Average	0.11021	0.13265
	Stacked Models	0.10989	0.13169
		0.10959	0.13156
Final Submission	0.7 Stacked + 0.3 XGBoost	N/A	0.12873

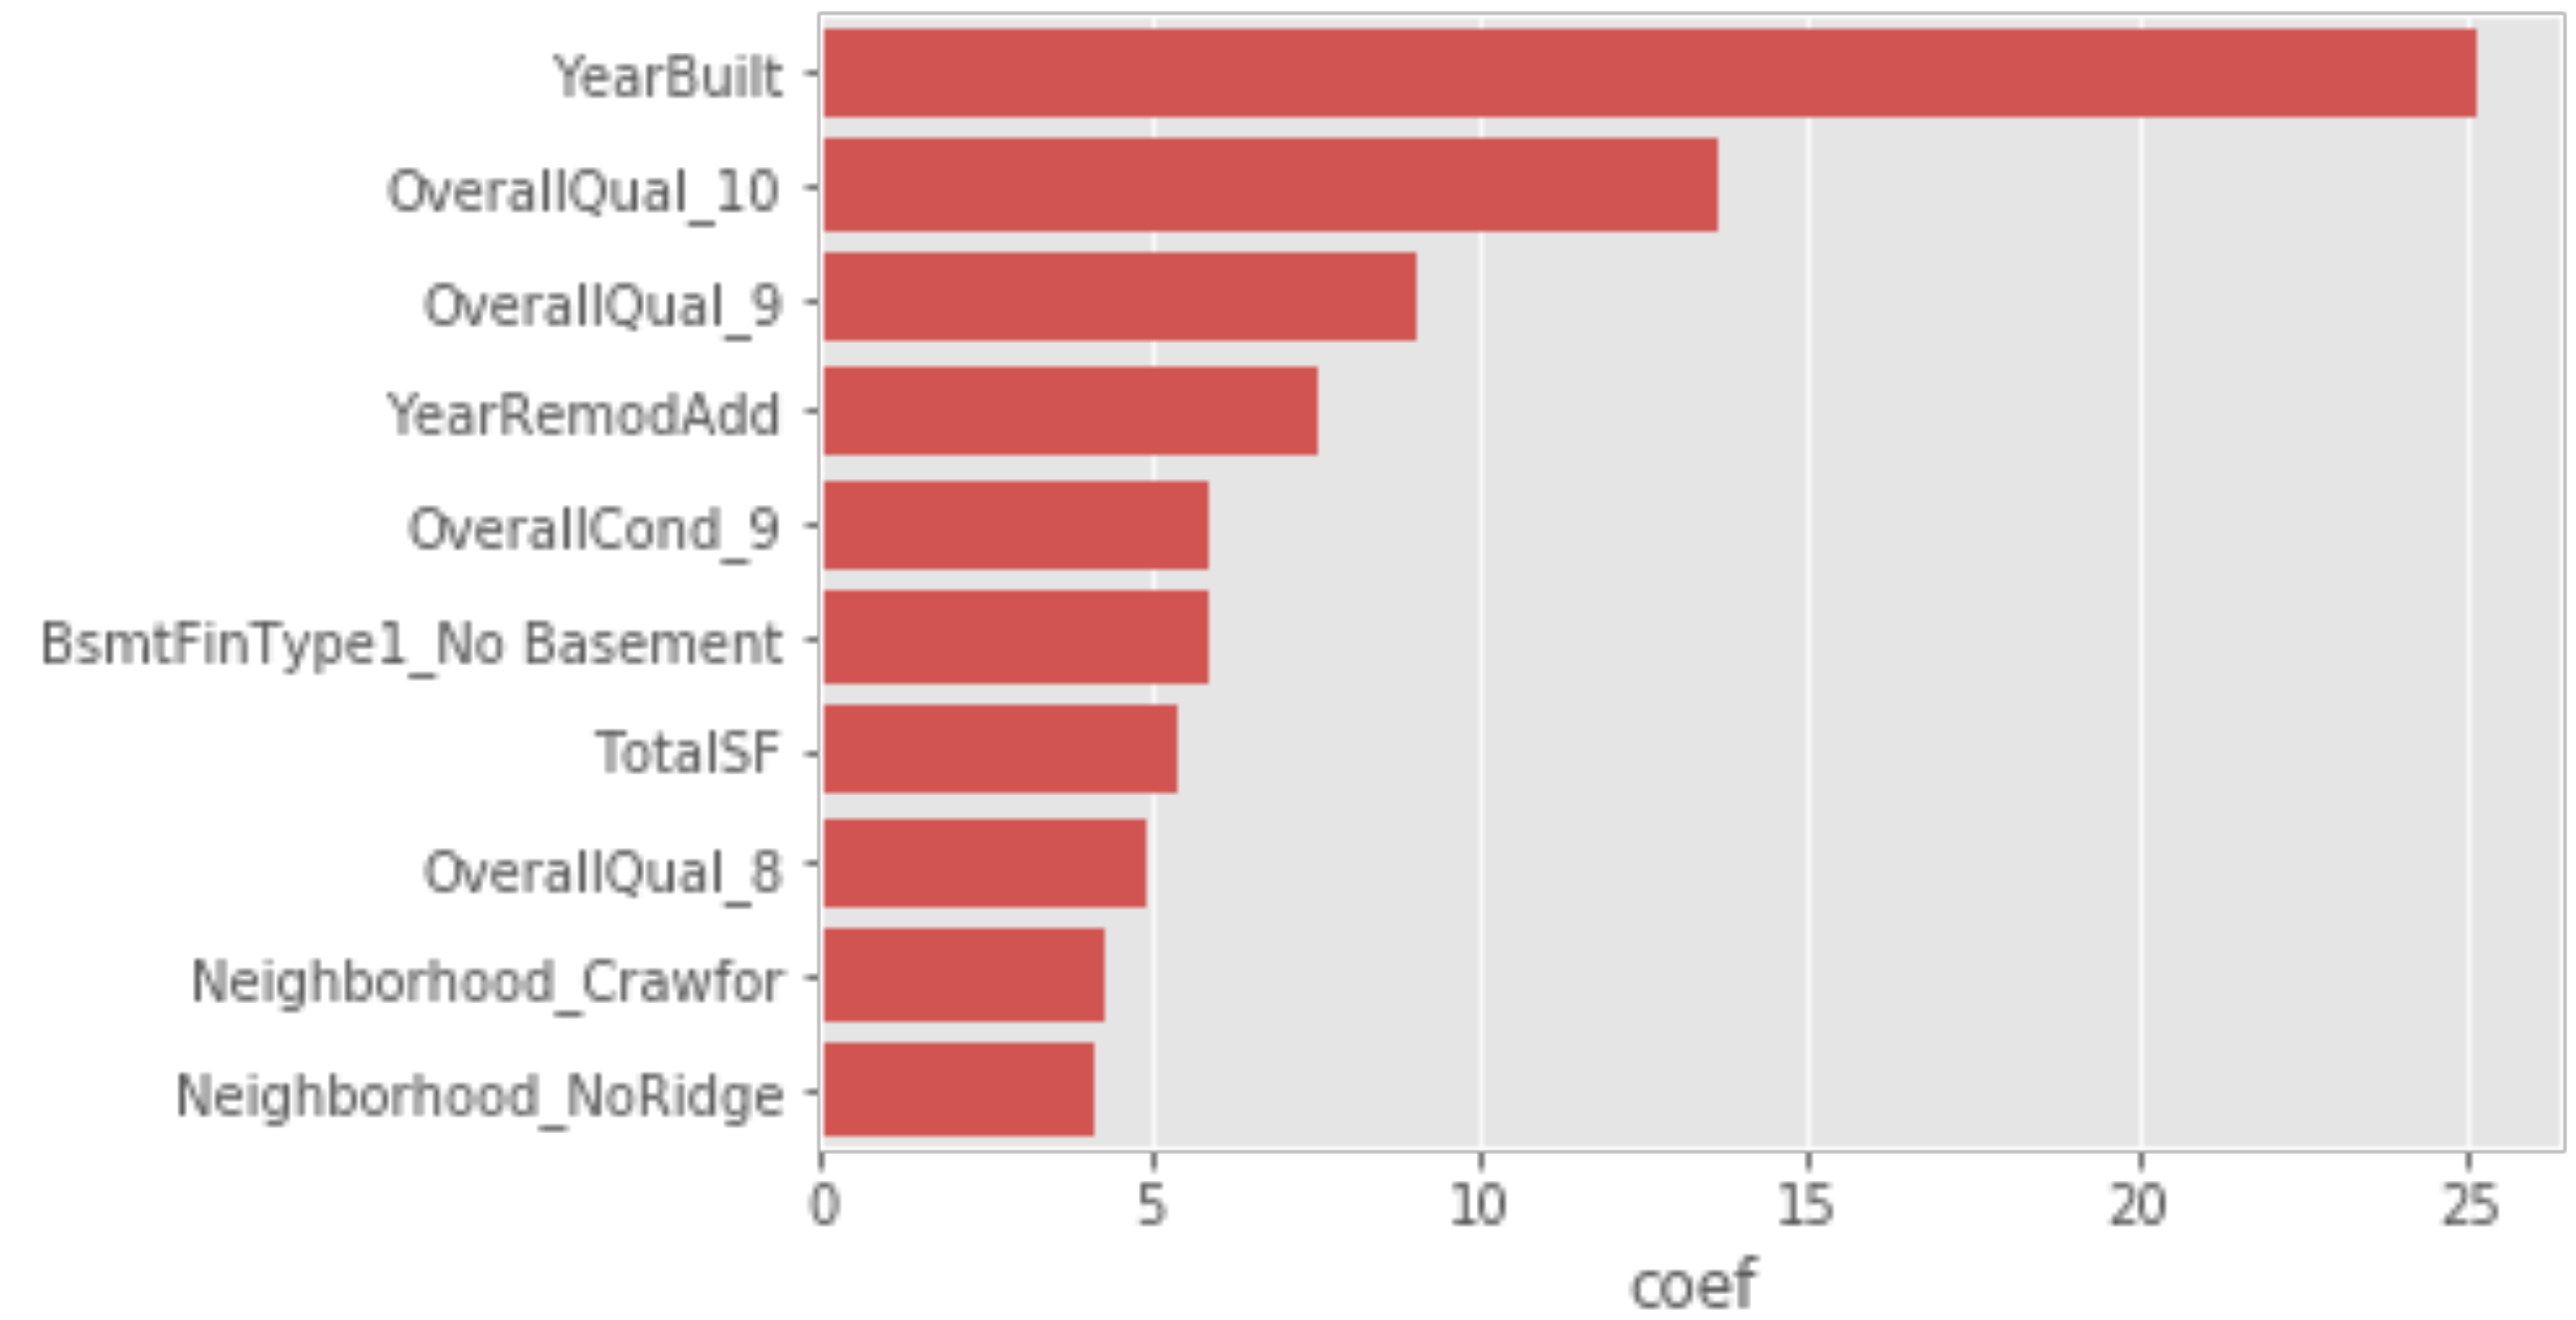

Model 1: Ridge





Model 2: Lasso





Model 3: Elastic Net



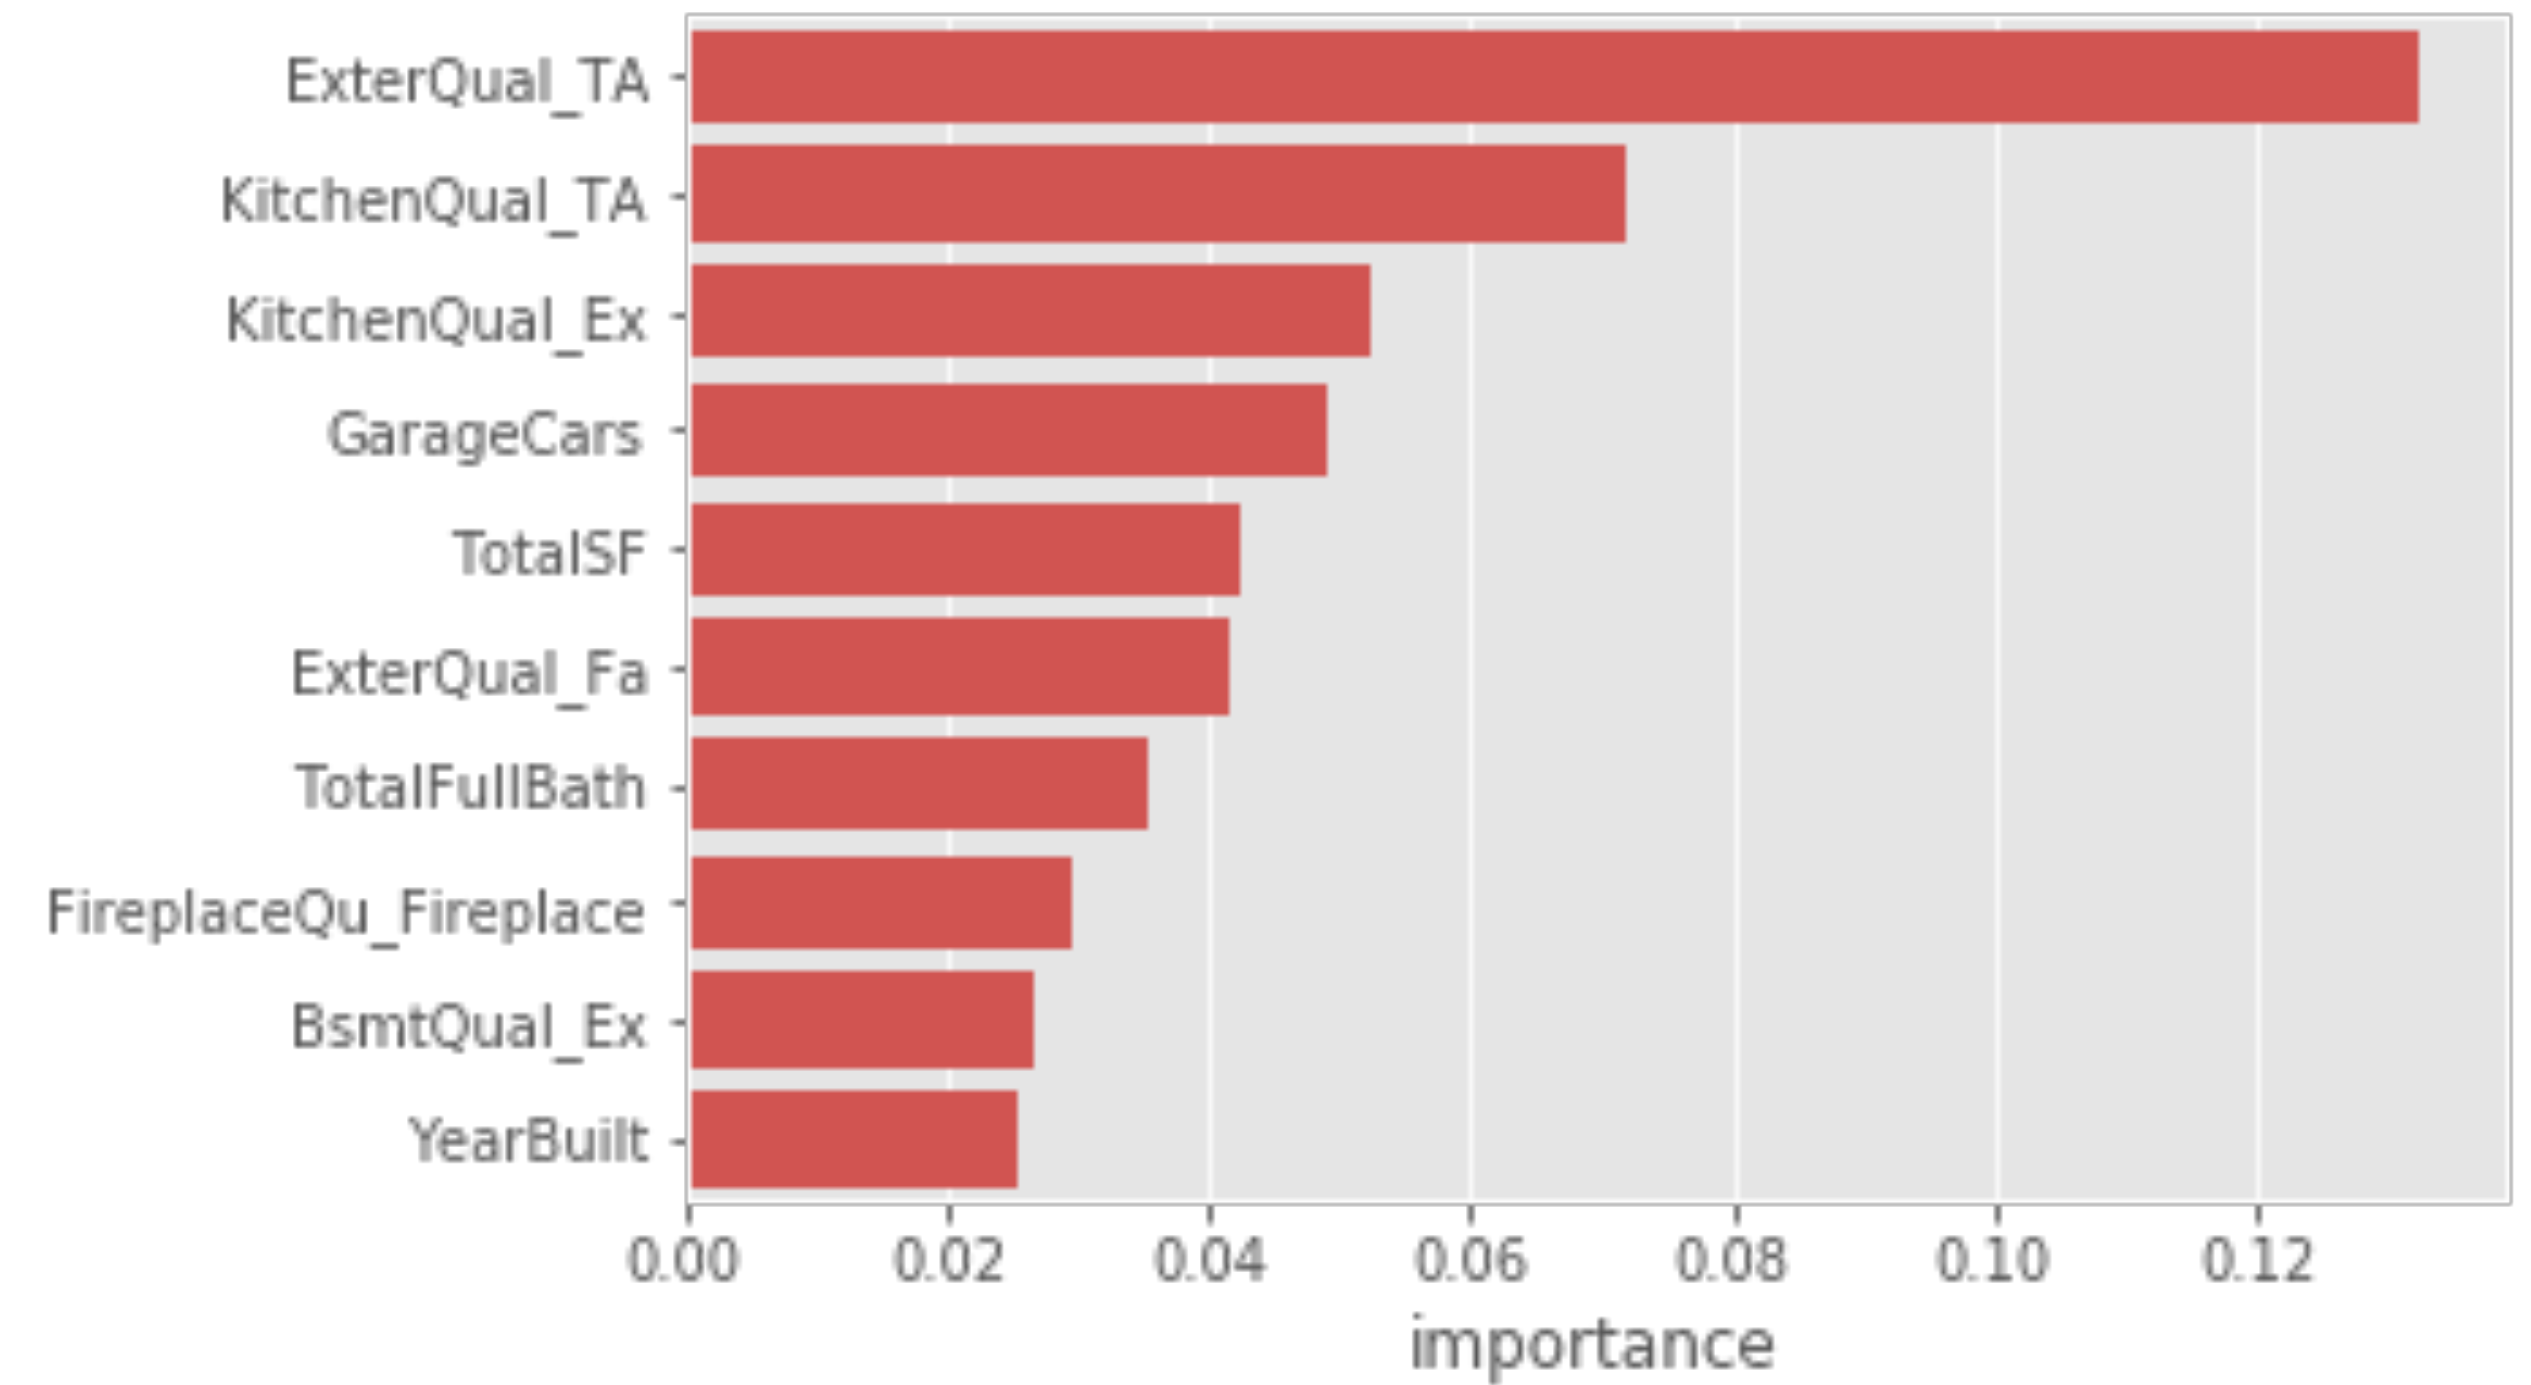

Model 4: Catboost





Model 5: Gradient Boost



Model 6: XGBoost



Model 7: LightGBM





Conclusion

Recommendation:

- Our best predictions of home portfolio value came from a weighted average of our Stacked and Light GBM models

Model	Details	Kaggle Score
Weighted Average of Best Models	0.7 Stacked Model + 0.3 Light GBM	0.12873

- This model performance is based on the Kaggle Score, which is the RMSE (root mean square error) of the log of the predictions and the log of the actual sales prices

Questions?

Thank You!